

Automated Hardware Trojan Detection in FPGAs

Nicholas Houghton, Samer Moein, and Fayez Gebali

Department of Electrical and Computer Engineering

University of Victoria

P.O. Box 1700 STN CSC

Victoria, B.C. V8W 2Y2

Email: {nhoughto, samerm, fayez}@uvic.ca

Abstract—Electronics have become such a staple in modern life that we are just as affected by their vulnerabilities as they are. Ensuring that the processors that control them are secure is paramount to our intellectual safety, our financial safety, our privacy, and even our personal safety. The market for integrated circuits is steadily being consumed by a reconfigurable type of processor known as a field-programmable gate-array (FPGA). The very features that make this type of device so successful also make them susceptible to attack. FPGAs are reconfigured by software; this makes it easy for attackers to make modification. Such modifications are known as hardware trojans. There have been many techniques and strategies to ensure that these devices are free from trojans but few have taken advantage of the central feature of these devices. The configuration Bitstream is the binary file which programs these devices. By extracting and analyzing it, a much more accurate and efficient means of detecting trojans can be achieved. This discussion presents a new methodology for exploiting the power of the configuration Bitstream to detect and described hardware trojans. A software application is developed that automates this methodology for *Xilinx* FPGAs.

I. INTRODUCTION

The term *Trojan Horse* or *Trojan* has become a modern metaphor for a deception where by an unsuspecting victim welcomes a foe into an otherwise safe environment [1]. Since the dawn of the computer we have dealt with software threats. We are almost as good at protecting ourselves against them as attackers are at making them. In recent years a new incarnation of electronic danger has emerged; in hardware. In this new arena of attack and defend those who seek to defend are far behind.

Integrated Circuit (IC) designs for Field Programmable Gate-Arrays (FPGAs) are made using a software language known as a Hardware Description Language (HDL). The design is then converted to a binary file called a configuration Bitstream which is then downloaded onto the device; this process is known as synthesizing the design. There have been many attempts to develop mechanisms and techniques to determine whether a malicious user has tampered with the design via test vectoring or side-channel analysis. As of yet there has been little effort to directly analyze the configuration Bitstream.

A method of extracting and analyzing the configuration Bitstream to determine the presence of hardware trojans has been developed. This method is able to meaningfully read the long binary file and extract modifications. Any discovered changes are located on the device using a new technique referred to as 'Component Mapping'. Further, these changes are then mapped to the user's original design. Knowing

which components of the device have been modified, and the instances of the synthesized design allows for a powerful description to be built. A software tool known as FPGA Trojan Detector which implements this new method has been built. FPGA Trojan Detector is able to automatically detect and analyze trojans in *Xilinx* FPGAs. Once analyzed a meaningful description is provided using the trojan taxonomy presented in [2].

The contributions of this paper are:

- 1) A new method mapping configuration Bitstream words to device components named 'Component Mapping'.
- 2) A systematic process of detecting and analyzing hardware trojans in FPGAs
- 3) A software tool which automates these new methods.

II. METHODOLOGY

Figure 1 provides a visual representation of the use-case assumed for the purposes of this work. With the exception of the fabrication process, all stages of production of an FPGA implementation are assumed to have been done "in-house". Any trojan discovered is inserted in the fabrication phase; all other stages are trusted. The method of automated trojan detection described in this work would take place in the 'testing' phase of the life-cycle. Figure 2 shows an

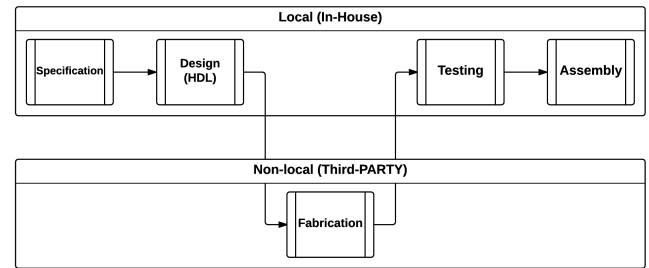


Fig. 1: FPGA Life-Cycle

overview of the trojan detection methodology. As mentioned, FPGA designs are written in a Hardware Description Language (HDL). *Xilinx* provides a series of User-Interface (UI) and command line tools to process the HDL known as the 'tool-chain'. The tool chain generates a series of files that are used for a variety of purposes as shown in the 'Resultant Files' box in Figure 2. The NGC file is a non-human readable semantic description of the design known as a netlist. This file

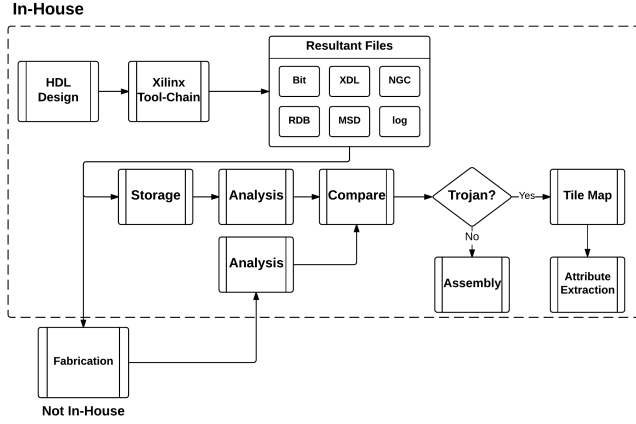


Fig. 2: Methodology Overview

can be converted into a human-readable version known as a *Xilinx* Design Language (XDL) file which will be used to relate any discovered modifications to the user's design. The Bit file is the binary representation of the design implementation. It is referred to as the Bitstream or 'configuration' Bitstream and is the final form that is loaded into the FPGA. The resultant files, produced 'in-house' are to be kept in secure storage while a copy is sent to be fabricated; these stored copies are referred to as Golden and assumed to be trojan-free. Though it is known that the fabrication houses will often attempt to make optimizations on designs, this methodology requires that no such efforts be made. When the completed batch of fabricated chips are returned the Bitstream is extracted from a sample using the *Xilinx* feature Readback. That which is extracted is referred to as the Target Bitstream. The Golden and Target Bitstreams are analyzed in conjunction to detect differences. Any discovered differences are then attributed to the corresponding component in the architecture, described in section III-A. Modified components are then related to the user's design using the XDL file. Finally, the resultant taxonomic description is built and returned to the user.

III. BITSTREAM AND USER DESIGN ANALYSIS

An FPGA is composed of a matrix of blocks referred to as the gate-array. Blocks perform functions specified by the configuration they receive from the synthesis process. The *Xilinx* Bitstream is a binary file composed of a series of 32-bit words organized into 'frames'. A frame is a string of single bits that span from the top to the bottom of a clock region of a device. A frame affects every block in a column and multiple horizontally adjacent frames are required to configure an entire column. Each frame is uniquely identified by a 32-bit address shown in Table I and is the smallest addressable element.

A. Component Mapping

The FPGA Trojan Detector employs a method referred to as Component Mapping to create a link between each word in a configuration frame and the component on the device that it configures. This information is not publicly released by *Xilinx* as a means of providing security through obscurity.

When powered-on the Bitstream is transmitted in frame address order to populate the dynamic memory in the tiles of the gate-array. The frame addressing scheme describes where in the gate-array the frame is destined fairly directly.

The Block Address (BA) value discerns whether the column type is a logic block column (BA: 0), a Block RAM (BRAM) column (BA: 1), or a memory interconnect block (BA: 2). Frames with a BA value of 1 are clearly destined to configure the BRAM and do not need further analysis for the purposes of this method. Frames with a BA value of 0 or 2 must be mapped more finitely. The row address specifies which row of clock-regions the frame is destined, as seen in Figure 3. As an example the Virtex-5 240T has 12 rows [3];

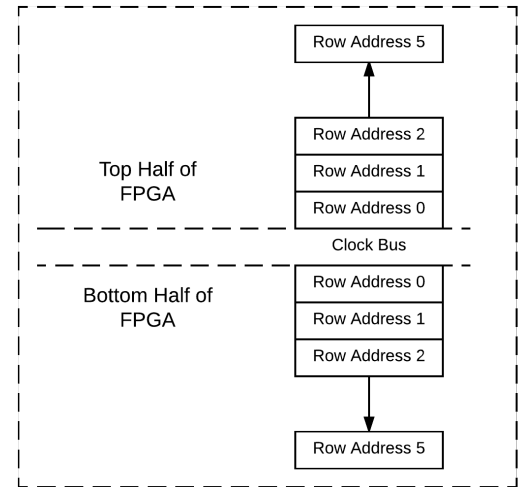


Fig. 3: Row Order of Virtex-5 Clock Region

its row address spans from 0-5 and the Top bit in the address indicates whether it is in the top or bottom half of the device in accordance with Figure 3. Once the correct clock region is discerned the major address is used to determine which column the frame configures. The major address begins at 0 on the left and counts up towards the number of columns in the row. Finally, the minor address is used to determine which sub-column has been modified.

In the *Xilinx* jargon, a tile is a conceptual encapsulation of components in the gate-array which perform a specific task. A block may contain multiple tiles. As an example, in a

TABLE I: Frame Address

Unused								BA			T	Row Address						Major Address							Minor Address							
31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
0	0	0	0	0	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	x	0	0	0	0	0	0	0	0	0	0

Configurable Logic Block (CLB) column, each block consists of an interconnect tile (also known as a Switching Matrix (SM)), an interface tile, and a CLB tile. The Frames which configure a column have a minor address numbered from left to right, starting with 0. Frames with a minor address from 0 to 25 configure the interconnect (SM). Frames numbered 26 and 27 access the interface tile and the remainder configure the CLB tile [3].

A frame is composed of 41 words that can be thought of as a vertical stack that aligns within a column. A column consists of a stack of blocks; there are 20 blocks per CLB column, 40 per IOB column, 4 per BRAM...etc for Virtex-5 devices. Equation 1 can be used to compute the number of 32-bit words that span each block.

$$n = (W - C) + B \quad (1)$$

where:

n = Number of Words per Block
 W = Number of 32-bit words per frame
 C = Number of clock words per frame
 B = Number of blocks per column

The 32-bit Words are addressed from the 'top' of a device down. Equation 2 can be used to map a particular word in a frame to a block in the column.

$$i = B - \left\lfloor \frac{w}{n} \right\rfloor \quad (2)$$

where:

i = Word Number in frame
 B = Number of blocks per column
 w = Word number
 n = Number of Words per Block

By using equations 1 and 2 and knowing how the minor address relates to the column it is possible to link any modifications in the Bitstream to its corresponding tile in the device.

B. Design Analysis

IV. DETERMINING TROJAN ATTRIBUTES

The complexity of Integrated Circuit designs and their corresponding trojans requires a more human-friendly scope. The taxonomy in [2] provides a series of thirty-three attributes which a trojan may or may not possess. Though it is desirable to be able to observe and directly extract the attributes a trojan possesses it is not always possible. In [2] a matrix, \mathbf{R} , is provided which describes the relationships between each of the thirty-three attributes in the taxonomy. When it is not possible to directly determine the presence of certain attributes, this relation matrix is used to infer their existence. The analysis stage of the automated trojan detection technique provided by this work begins by extracting those attributes that are directly observable then using matrix \mathbf{R} to infer the existence of the remainder.

A. Observed Location Attributes

The presence of attributes in the *Location* category are directly observable from the results of the component mapping method described in section III-A. *Xilinx* tiles conform to purpose-specific groups or block types which were discussed in section III. These block types contain sub-types that perform actions which pertain to the *Location*, category.

- 1) The **Processor** attribute pertains to the core functionality of the design logic. It can be awarded for presence of a modified CLB tile or Interconnect tile.
- 2) The **Memory** attribute can be awarded for the presence of modified BRAM components.
- 3) The **IO** attribute can be awarded for presence of modified IOB tiles.
- 4) The **Power Supply** attribute can be awarded for the presence of modified interface or configuration tiles.
- 5) The **Clock Grid** attribute can be awarded for modified clock tiles.

B. Scatter Score Method

The gate-array configuration of components in *Xilinx* FPGAs allows for an analytical method of determining attributes in the *Physical Layout* category. The "Scatter Score" method uses the grid coordinates of components to derive a numerical score rating for the size, position, and augmentation of configured tiles. Tiles are assigned global coordinates that represent their horizontal and vertical positions within the gate array denoted x and y respectively. These values can then be used to strongly infer the presence of *Physical Location* attributes.

The golden chip is first analyzed. The set of all tiles which are configured in the golden design is found and a series of numerical descriptors are computed.

$$n = \sum_{x=0}^X \sum_{y=0}^Y T_{xy} \quad (3)$$

where:

n = Number of all **configured** tiles
 X = The column width of the gate-array
 Y = The number of rows of the gate-array
 T = A configured tile

$$a_x = \frac{1}{n} \sum_{x=0}^n T_x \quad (4) \quad a_y = \frac{1}{n} \sum_{y=0}^n T_y \quad (5)$$

$$\sigma_x = \sqrt{\frac{1}{n} \sum_{x=0}^X (x_i - a_x)^2} \quad \sigma_y = \sqrt{\frac{1}{n} \sum_{y=0}^Y (y_i - a_y)^2} \quad (6) \quad (7)$$

where:

a_x = The average x coordinate of configured tiles
 a_y = The average y coordinate of configured tiles
 T_x = The x coordinate of a configured tile
 T_y = The y coordinate of a configured tile
 σ_x = The standard deviation of the x coordinate of configured tiles
 σ_y = The standard deviation of the y coordinate of configured tiles

Equations 4 and 5 are used to create a rating known as the Position Median in Equation 8. The Position Median value

provides a simple descriptor for where in the gate array the design is centralized. The Scatter Score in Equation 9 describes how spread out or, *clustered* the design is.

$$M_{xy} = (a_x, a_y) \quad (8) \quad S_{xy} = (\sigma_x, \sigma_y) \quad (9)$$

where:

M_{xy} = The Position Median

S_{xy} = The Scatter Score

The results of the component mapping method described in section ?? are used to generate the set of all tiles reconfigured by the trojan. The set of reconfigured tiles can be said to contain three subsets: the subset of tiles activated by the trojan, those deactivated and those modified. The results of the golden design analysis, the subsets, and the numeric descriptors can be used to discern which of the *Physical Location* attributes the trojan possesses. The *Physical Location* category contains six attributes. These six can be considered three pairs; a trojan exhibits one attribute from each pair.

- 1) **Large or Small** (attributes 23 or 24): According to [2], small trojans are defined as those that are nearly impossible to detect via power consumption. From this it can be said that 'small' trojans occupy minimal resources. Trojans where the number of reconfigured tiles is less than 5% of the number of tiles in the golden design are considered small. Other wise they are attributed as large.
- 2) **Changed Layout or Augmented** (attributes 25 or 26): A 'changed layout' trojan is such that only tiles that are configured by the golden design are reconfigured. An augmented trojan is where additional layout is added. The presence of 'activated' or 'deactivated' tiles indicates an augmented trojan.
- 3) **Clustered or Distributed** (attributes 27 or 28): The trojan is considered to be clustered when the standard deviation of the reconfigured tile positions is less than 15%; distributed otherwise.

C. Insertion and Abstraction Attributes

The linear nature of the manufacturing life-cycle implies a propagation of effects. For the purposes of this method it is assumed that the only non-trustworthy stage in the life-cycle is fabrication. In other words, the trojan was inserted in the third-party fabrication stage. Due to the propagating nature of the life-cycle the effects of the modifications made in the fabrication stage (attribute 3) are felt in the testing (attribute 4) and assembly (attribute 5) stages. Hence, it can be said that this trojan possesses insertion category attributes 3, 4 and 5. FPGA designs are made with a HDL. These languages dictate component arrangement in the Registry Transfer Level (RTL) abstraction level. Hence it can be said that trojans occurring in FPGAs take place in the System (attribute 6) and RTL (attribute 7).

D. Relation Matrix Use

Attributes which are not directly observable can be inferred using a systematic method of analyzing the rows and columns of the relation matrix presented in [2]. The FPGA Trojan Detector takes the attributes it is able to directly observe and uses them as input to this process. A thorough description of this method is given in [4] and [5].

V. FPGA TROJAN DETECTOR

The FPGA Trojan Detector can provide manufacturers additional security that takes only minutes to complete; the analysis requires only a few button clicks. It was built to be a stand-alone application; it is light-weight, cross platform and does not require installation.

A. Technologies Used

1) *Xilinx* : *Xilinx* is one of the two largest manufacturers of FPGAs; their devices are considerably more popular than their competitors. FPGA Trojan Detector takes

2) *Java*: FPGA Trojan Detector was written in Java primarily in order to interface with the API known as *RapidSmith* which is described in section V-A3. However, Java additionally allows the FPGA Trojan Detector to be a compact, cross-platform application. The custodians of the Java language, Sun Microsystems, promote the slogan 'write-once, run anywhere'. The Java language provides a native Graphical User-Interface (GUI) toolkit known as *Swing*. *Swing* is an API that is part of Oracle's Java Foundation Classes; in other words it is readily available to all users of Java.

3) *RapidSmith* : *RapidSmith* is an Application Programming Interfaces (API) written in Java that enables Computer Aided Design (CAD) tool creation for *Xilinx* FPGAs [6]. Its purpose is to be used as a rapid prototyping platform for experimentation and research. It was chosen as a supporting library for the FPGA Trojan Detector for several reasons. First, the code base provides a series of class structures that astutely mirror the architecture of *Xilinx* devices. Secondly, it provides ready-made tools for extracting the configuration frames from Bitstream files. Bitstream files are long binary sequences, without the tools provided by *RapidSmith* the analysis of these files becomes an arduous task. Finally, and most importantly, the creators of *RapidSmith* have developed a means of condensing XDLRC files into a greatly compressed format referred to as a 'database' file. The FPGA Trojan Detector requires considerable detail of an FPGA's architecture in order to accurately described the effects a trojan has on a design.

VI. RESULTS

To demonstrate its potential, FPGA Trojan Detector has been tested using a series of benchmarks. These benchmarks are included in the FPGA Trojan Detector application package as an example for users.

A. Priority Decoder

The FPGA Trojan Detector was first tested using a small 'Priority Decoder' circuit presented by F. Brglez and H. Fujiwara in [7]. The provided verilog code for the decoder benchmark was synthesized on a Virtex-5 XC5VLX155 and the generated Bitstream and XDL files were acquired. The priority decoder Bitstream file was fed to both the Golden and Target inputs to the FPGA Trojan Detector. Feeding the same Bitstream file to both inputs replicates the occurrence where the third-party fabrication house made no modifications; the Bitstream extracted from the Target device is exactly the same as the Golden. It is expected that the FPGA Trojan

Detector returns a result indicating that there is no trojan present. The FPGA Trojan Detector successfully analyzed the Bitstreams and determined that there was no trojan present, as expected.

B. User Authentication Circuit

Consider a circuit designed to compute a function $F(x)$ for a system to authenticate user-password pairs x and $F(x)$. The system performs the arithmetic operation $F(x) = x^2$ to validate users. The customer wishes to provide access to ten users labeled I_0 to I_9 . To identify all ten users, four input bits are required, x_1 to x_4 . The largest function output is 81 meaning seven bits are required for output, Z_1 to Z_7 , as illustrated in Fig. 4. A trojan can be inserted into this circuit as shown in Fig. 5 (called a backdoor trojan). The




Fig. 4: A Simple User Authentication Circuit




Fig. 5: The Back-door Trojan

TABLE II: Outputs of the Circuits in Figs. 4 and 5 [2]

	User Id	Input	Circuit A	Circuit B
Inputs	I_0	0	0	0
	I_1	1	1	1
	I_2	2	4	4
	\vdots	\vdots	\vdots	\vdots
	I_9	9	81	81
DC	I_{10}	10	68	100
	I_{11}	11	89	121
	\vdots	\vdots	\vdots	\vdots

outputs of the original and infected circuits are compared in Table II. A simple test will show that the circuit outputs the desired $F(x) = x^2$ for each of the users. However, upon closer

inspection it is noted that the inputs corresponding to $x = 10$ to $x = 15$ are not used; there are no clients occupying those identifications. These unused inputs are referred to as 'don't-cares' (DC), meaning that it is not important to the function of the circuit what their corresponding output is. Don't-care cases are a typical vulnerability which can be exploited by an attacker.

The circuits shown in Figures 4 and 5 were implemented and synthesized on a Virtex-5 240T (XC5VSX240T). The resultant Bitstreams were input to the FPGA Trojan Detector and the system output the attributes:

- Attribute 3: Fabrication
- Attribute 4: Testing
- Attribute 5: Assembly
- Attribute 6: System
- Attribute 7: RTL
- Attribute 12: Change in Functionality
- Attribute 17: Combinational
- Attribute 18: Functional
- Attribute 20: Always On
- Attribute 24: Small
- Attribute 26: Augmented
- Attribute 27: Clustered
- Attribute 29: Processor

As expected the results state that the trojan was inserted in the Fabrication phase (3); it is the earliest stage in the 'insertion' phase produced indicating it as the source. The effects of this modification propagate to the Testing (4) and Assembly (5) phases as expected. The modifications reach both the System (6) and Registry Transfer Level (RTL) (7) abstraction levels. Since the modifications were made using the schematic designer provided by Xilinx which works in the RTL level, these results are as expected. The results indicate that the trojan Changes Functionality (12). This agrees with the modification to the values listed in Table II. The trojan does not take affect over multiple clock cycles; this indicates it is composed of only Combinational (17) circuitry which is reflected in the results. The trojan did not modify power levels or operation configurations, only design configurations. This indicates that the trojan can be described as Functional (18), not Parametric (19); this agrees with the results. Since the modification made a permanent alteration to the internal wiring of the circuit it can be said to be Always On (20). The modified values for input $x = 10$ or $x = 11$ are always available and not activated. This is consistent with the returned results. The trojan changed only minor routing configurations in the circuit designs, this required the alteration of only a few tiles. The new route required the activation of tiles which were not active in the Golden design; further, all of the modified tiles are nearby those in the Golden design. With these observations it can be said that this trojan exhibits Physical Layout attributes Small (24), Augmented (26) and Clustered (28). All of the expected Physical Layout attributes were correctly determined by the Scatter Score method of section IV-B. Finally, all of the tiles modified by the trojan belong to major block type 0: Logic Type. These tiles only affect the internal processing of the circuit. No IOB, Clock or BRAM tiles were modified. This is reflected by the fact that only Location attribute, Processor (29), was returned by the analysis. The results observed by the experiment conformed with the experiments expectations

demonstrating the accuracy of the method. The entire analysis takes less than a minute to perform.

C. AES-T100

In 2013 H. Salmani, M. Tehranipoor and R. Karri published a discussion on the design and development of FPGA trojan benchmarks. They collaboratively developed a series of verilog, VHDL and virtual machines that demonstrated effective creation of testable benchmarks. They took the benchmarks they created and published them on a website they created called *Trust-Hub*. To demonstrate the efficacy of the FPGA Trojan Detector a benchmark named 'AES-T100' was chosen. From the description it is reasonable to expect certain results from the FPGA Trojan Detector. The supporting documentation claims that the trojan 'leaks the secret key'. From this we should expect our results to contain Effect attribute Information Leakage (13). Information being leaked from a device will need a means to be transmitted to the attacker. Location attribute IO (31) may be observed. It then states that it leaks 'single bits over many clock cycles'. This suggests that the trojan exhibits some form of Sequential Logic (16). This may or may not require modification to clock tiles; Location attribute Clock Grid (33) may be observed. It then states that the PRNG is initialized to a predefined value; initialization requires the value be stored in memory. Location attribute Memory (30) should be expected. The trojan then uses a 'power side-channel' as a communication channel. This will require modification to power tiles; Location attribute Power Supply (32) should be expected.

The source code for the Golden and Target designs were downloaded from *trust-hub.org* and synthesized on a Virtex-5 240T (XC5VSX240T). The resultant Bitstreams were analyzed and the FPGA Trojan Detector output the following attributes which correlate well to the description in the documentation:

- Attribute 3: Fabrication
- Attribute 4: Testing
- Attribute 5: Assembly
- Attribute 6: System
- Attribute 7: RTL
- Attribute 13: Information Leakage
- Attribute 16: Sequential
- Attribute 18: Functional
- Attribute 20: Always On
- Attribute 24: Large
- Attribute 26: Augmented
- Attribute 27: Distributed
- Attribute 29: Processor
- Attribute 30: Memory
- Attribute 31: IO
- Attribute 32: Power Supply
- Attribute 33: Clock Grid

VII. CONCLUSION

Configuration Bitstreams are enormous strings of binary data. To the human reader this information means nothing. To an FPGA, however, this data is everything. Every conceivable design, and every possible trojan is contained within the Bitstream. Yet, due to the sheer volume of information within it and the lack of details on its format it has not previously

been a common subject of study. With its success, Integrated Circuit manufacturers that use FPGAs will have an additional tool to ensure that their products operate as expected. Using the FPGA Trojan Detector takes only a few button clicks on the User-Interface. Its simple construction does not require any additional software or complicated install procedures and it can be used on any major operating system. Ensuring chips that have returned from Fabrication operate as expected takes no more than a few minutes. With the use of the FPGA Trojan Detector manufacturers will not need to train employees, buy expensive equipment or waste man-hours on additional testing.

REFERENCES

- [1] Michael Wood. *"In search of the Trojan war.* 1st ed. The British Broadcasting Corporation, 1998. ISBN: 978-0-520-21599-3.
- [2] S. Moein et al. "An attribute based classification of hardware trojans". In: *Computer Engineering Systems (ICCES), 2015 Tenth International Conference on.* 2015, pp. 351–356. DOI: 10.1109/ICCES.2015.7393074.
- [3] *Virtex-5 FPGA Configuration User Guide.* v3.11. Xilinx. 2012.
- [4] Samer Moein. "Systematic Analysis and Methodologies for Hardware Security". PhD thesis. 3800 Finnerty Rd, Victoria, BC, Canada V8P 5C2: University of Victoria, Aug. 2015.
- [5] N. Houghton et al. "An Automated Web Tool for Hardware Trojan Classification". In: *ESC'16 2* (July 2015).
- [6] C. Lavin et al. "RapidSmith: Do-It-Yourself CAD Tools for Xilinx FPGAs". In: *2011 21st International Conference on Field Programmable Logic and Applications.* 2011, pp. 349–355. DOI: 10.1109/FPL.2011.69.
- [7] F. Brglez, D. Bryan, and K. Kozminski. "Combinational profiles of sequential benchmark circuits". In: *Circuits and Systems, 1989., IEEE International Symposium on.* 1989, 1929–1934 vol.3. DOI: 10.1109/ISCAS.1989.100747.